

```
import pandas as pd

# File paths
file_paths = {
    "2022-2023": "/content/gsw-shg-en 2022-2023.csv",
    "2023-2024": "/content/gsw-shg-en 2023-2024.csv",
    "2024-2025": "/content/gsw-shg-en 2024-2025.csv",
}

# Load datasets
dataframes = {year: pd.read_csv(path) for year, path in file_paths.items()}

# Display basic information about the datasets
{year: df.info() for year, df in dataframes.items()}
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 207192 entries, 0 to 207191
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   grain_week            207192 non-null int64
1   crop_year             207192 non-null object
2   week_ending_date      207192 non-null object
3   worksheet             207192 non-null object
4   metric                207192 non-null object
5   period               207192 non-null object
6   grain                207192 non-null object
7   grade                62130 non-null object
8   region               202176 non-null object
9   Ktonnes              207192 non-null object
dtypes: int64(1), object(9)
memory usage: 15.8+ MB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 211116 entries, 0 to 211115
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   grain_week            211116 non-null int64
1   crop_year             211116 non-null object
2   week_ending_date      211116 non-null object
3   worksheet             211116 non-null object
4   metric                211116 non-null object
5   period               211116 non-null object
6   grain                211116 non-null object
7   grade                61092 non-null object
8   region               205404 non-null object
9   Ktonnes              211114 non-null object
dtypes: int64(1), object(9)
memory usage: 16.1+ MB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 98581 entries, 0 to 98580
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   grain_week            98581 non-null int64
1   crop_year             98581 non-null object
2   week_ending_date      98581 non-null object
3   worksheet             98581 non-null object
4   metric                98581 non-null object
5   period               98581 non-null object
6   grain                98581 non-null object
7   grade                28878 non-null object
8   region               95893 non-null object
9   Ktonnes              98581 non-null object
dtypes: int64(1), object(9)
memory usage: 7.5+ MB
{'2022-2023': None, '2023-2024': None, '2024-2025': None}
```

```
# Convert data types and handle missing values
```

```
for year, df in dataframes.items():
    # Convert week_ending_date to datetime
```

```


df["week_ending_date"] = pd.to_datetime(df["week_ending_date"], errors='coerce')

# Convert Ktonnes to numeric, forcing errors to NaN (to clean non-numeric values)
df["Ktonnes"] = pd.to_numeric(df["Ktonnes"], errors='coerce')

# Fill missing region values with 'Unknown' for now
df["region"].fillna("Unknown", inplace=True)

# Verify the changes
dataframes["2022-2023"].info(), dataframes["2023-2024"].info(), dataframes["2024-2025"].info()

```



#	Column	Non-Null Count	Dtype
0	grain_week	207192 non-null	int64
1	crop_year	207192 non-null	object
2	week_ending_date	85134 non-null	datetime64[ns]
3	worksheet	207192 non-null	object
4	metric	207192 non-null	object
5	period	207192 non-null	object
6	grain	207192 non-null	object
7	grade	62130 non-null	object
8	region	207192 non-null	object
9	Ktonnes	204464 non-null	float64

dtypes: datetime64[ns](1), float64(1), int64(1), object(7)
memory usage: 15.8+ MB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 211116 entries, 0 to 211115
Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	grain_week	211116 non-null	int64
1	crop_year	211116 non-null	object
2	week_ending_date	82160 non-null	datetime64[ns]
3	worksheet	211116 non-null	object
4	metric	211116 non-null	object
5	period	211116 non-null	object
6	grain	211116 non-null	object
7	grade	61092 non-null	object
8	region	211116 non-null	object
9	Ktonnes	208670 non-null	float64

dtypes: datetime64[ns](1), float64(1), int64(1), object(7)
memory usage: 16.1+ MB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 98581 entries, 0 to 98580
Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	grain_week	98581 non-null	int64
1	crop_year	98581 non-null	object
2	week_ending_date	41153 non-null	datetime64[ns]
3	worksheet	98581 non-null	object
4	metric	98581 non-null	object
5	period	98581 non-null	object
6	grain	98581 non-null	object
7	grade	28878 non-null	object
8	region	98581 non-null	object
9	Ktonnes	97880 non-null	float64

dtypes: datetime64[ns](1), float64(1), int64(1), object(7)
memory usage: 7.5+ MB
<ipython-input-2-3f6985f1d881>:11: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[c

```

df["region"].fillna("Unknown", inplace=True)
(None, None, None)

```

```

from statsmodels.tsa.stattools import adfuller
import matplotlib.pyplot as plt
import seaborn as sns

```

```

# Filter data for Corn and Soybeans from Bay & Lakes and St. Lawrence regions
selected_regions = ["Bay & Lakes", "St. Lawrence"]

```

```
selected_regions = ["Soybeans", "Soybeans"]
selected_grains = ["Corn", "Soybeans"]

# Combine all years into one DataFrame for analysis
df_all_years = pd.concat(dataframes.values())

# Filter dataset for relevant grains and regions
df_filtered = df_all_years[df_all_years["grain"].isin(selected_grains) & df_all_years["region"].isin(selected_regions)]

# Aggregate weekly exports by summing up Ktonnes
df_grouped = df_filtered.groupby(["week_ending_date", "grain", "region"])["Ktonnes"].sum().reset_index()

# Pivot data for stationarity test
df_pivot = df_grouped.pivot(index="week_ending_date", columns=["grain", "region"], values="Ktonnes")

# Perform Augmented Dickey-Fuller (ADF) test for stationarity
adf_results = {}
for col in df_pivot.columns:
    adf_test = adfuller(df_pivot[col].dropna())
    adf_results[col] = {"ADF Statistic": adf_test[0], "p-value": adf_test[1]}

# Convert results to DataFrame for easier visualization
adf_results_df = pd.DataFrame(adf_results).T
adf_results_df
```




level_0	level_1	ADF Statistic	p-value
Corn	Bay & Lakes	-3.5306910497791852	0.007232992970737467
Corn	St. Lawrence	-2.199186693232346	0.20654077271486682
Soybeans	Bay & Lakes	-2.9613273676733143	0.038667186918730334
Soybeans	St. Lawrence	-8.394270486917113	2.3217472155750417e-13

Show 25 per page




Like what you see? Visit the [data table notebook](#) to learn more about interactive tables.


Distributions




2-d distributions




Values




Distributions




Categorical distributions



2-d distributions




Values



2-d categorical distributions






Faceted distributions


<string>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and




<string>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and




<string>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and



<string>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and



Next steps:

[Generate code with adf_results_df](#)

[View recommended plots](#)

[New interactive sheet](#)

```
from statsmodels.tsa.seasonal import STL
```

```
# Plot STL decomposition for each grain-region combination
```

```
fig, axes = plt.subplots(4, 1, figsize=(12, 16))
```

```
plt.suptitle("STL Decomposition of Weekly Grain Exports")
```

```
for i, col in enumerate(df_pivot.columns):
```

```
    series = df_pivot[col].dropna()
```

```
    stl = STL(series, period=52) # Weekly data, so seasonality period = 52 weeks
```

```
    result = stl.fit()
```

```
    # Plot trend component
```

```
    axes[i].plot(series.index, result.trend, label=f"{col} - Trend", color="blue")
```


```
    axes[i].plot(series.index, result.seasonal, label=f"{col} - Seasonality", color="orange", alpha=0.7)
```

```
    axes[i].set_title(f"STL Decomposition: {col}")
```

```
    axes[i].legend()
```

```
plt.tight_layout()
```

```
plt.show()
```




```
from statsmodels.tsa.seasonal import STL
import matplotlib.pyplot as plt

# Assuming df_pivot is your pivoted DataFrame with date index
fig, axes = plt.subplots(4, 1, figsize=(12, 16))
plt.suptitle("STL Decomposition of Weekly Grain Exports")

for i, col in enumerate(df_pivot.columns):
    series = df_pivot[col].dropna()
    stl = STL(series, period=52) # Weekly data, so seasonality period = 52 weeks
    result = stl.fit()

    # Plot trend component
    axes[i].plot(series.index, result.trend, label=f"{col} - Trend", color="blue")
    axes[i].plot(series.index, result.seasonal, label=f"{col} - Seasonality", color="orange", alpha=0.7)
    axes[i].set_title(f"STL Decomposition: {col}")
    axes[i].legend()

plt.tight_layout()
plt.show()
```




```
# Generate summary statistics for grain exports
summary_stats = df_filtered.groupby(["grain", "region"])["Ktonnes"].describe()

# Compute correlation matrix for numerical features
correlation_matrix = df_pivot.corr()

summary_stats, correlation_matrix
```

```
(
  grain    region    count    mean    std  min  25%  50%   75%  \
  Corn    Bay & Lakes  3220.0  61.309938  138.179232  0.0  0.6  6.9  46.025
         St. Lawrence  3272.0  27.109077   96.089543  0.0  0.0  0.5   9.300
  Soybeans Bay & Lakes  2947.0  78.695521  165.297619 -2.5  0.2  7.2  42.800
         St. Lawrence  2824.0  27.932011   91.097195  0.0  0.0  0.4   6.950

                                     max
  grain    region
  Corn    Bay & Lakes    998.5
         St. Lawrence   889.6
  Soybeans Bay & Lakes    986.0
         St. Lawrence   982.8
,
  grain    region    count    mean    std  min  25%  50%   75%  \
  region    region    Bay & Lakes  St. Lawrence  Bay & Lakes  St. Lawrence
  grain    region
  Corn    Bay & Lakes    1.000000    0.634776    0.523569   -0.300226
         St. Lawrence    0.634776    1.000000    0.228064   -0.265257
  Soybeans Bay & Lakes    0.523569    0.228064    1.000000    0.460453
         St. Lawrence   -0.300226   -0.265257    0.460453    1.000000)
```

```
# Generate summary statistics for grain exports
summary_stats = df_filtered.groupby(["grain", "region"])["Ktonnes"].describe()
print(summary_stats)
```

```
(
  grain    region    count    mean    std  min  25%  50%   75%  \
  Corn    Bay & Lakes  3220.0  61.309938  138.179232  0.0  0.6  6.9  46.025
         St. Lawrence  3272.0  27.109077   96.089543  0.0  0.0  0.5   9.300
  Soybeans Bay & Lakes  2947.0  78.695521  165.297619 -2.5  0.2  7.2  42.800
         St. Lawrence  2824.0  27.932011   91.097195  0.0  0.0  0.4   6.950

                                     max
  grain    region
  Corn    Bay & Lakes    998.5
         St. Lawrence   889.6
  Soybeans Bay & Lakes    986.0
         St. Lawrence   982.8
```

```
# Compute correlation matrix for numerical features
correlation_matrix = df_pivot.corr()
print(correlation_matrix)
```

```
(
  grain    region    count    mean    std  min  25%  50%   75%  \
  region    region    Bay & Lakes  St. Lawrence  Bay & Lakes  St. Lawrence
  grain    region
  Corn    Bay & Lakes    1.000000    0.634776    0.523569   -0.300226
         St. Lawrence    0.634776    1.000000    0.228064   -0.265257
  Soybeans Bay & Lakes    0.523569    0.228064    1.000000    0.460453
         St. Lawrence   -0.300226   -0.265257    0.460453    1.000000)
```

```
pip install ydata-profiling
```

 [Show hidden output](#)

```
import pandas as pd
from ydata_profiling import ProfileReport
```

```
# Load your cleaned dataset
df = pd.concat(dataframes.values()) # Assuming you combined all datasets
```

```
# Generate the profiling report
profile = ProfileReport(df, title="Grain Export EDA Report", explorative=True)

# Save the report as an HTML file
profile.to_file("EDA_Report.html")

print("EDA report generated: EDA_Report.html")
```

 [Upgrade to ydata-sdk](#)

Improve your data and profiling with ydata-sdk, featuring data quality scoring, redundancy detection, outlier identification, text validation, and synthetic data generation.

Summarize dataset: 100%

24/24 [00:11<00:00, 4.40it/s, Completed]

```
0%|          | 0/10 [00:00<?, ?it/s]
10%|█         | 1/10 [00:00<00:02, 3.67it/s]
20%|██        | 2/10 [00:02<00:10, 1.37s/it]
40%|████      | 4/10 [00:04<00:05, 1.02it/s]
50%|█████     | 5/10 [00:05<00:06, 1.22s/it]
60%|██████    | 6/10 [00:06<00:04, 1.10s/it]
70%|████████  | 7/10 [00:07<00:03, 1.15s/it]
100%|█████████| 10/10 [00:09<00:00, 1.10it/s]
```

```
/usr/local/lib/python3.11/dist-packages/ydata_profiling/model/correlations.py:87: UserWarning: There was an attempt to calculate correlations for a column with a non-numeric dtype. To hide this warning, disable the calculation
(using `df.profile_report(correlations={"auto": {"calculate": False}})`
If this is problematic for your use case, please report this as an issue:
https://github.com/ydataai/ydata-profiling/issues
(include the error message: 'cannot reindex on an axis with duplicate labels')
warnings.warn(
```

Generate report structure: 100%

1/1 [00:01<00:00, 1.37s/it]

Render HTML: 100%

1/1 [00:00<00:00, 2.35it/s]

Export report to file: 100%

1/1 [00:00<00:00, 93.97it/s]

EDA report generated: EDA_Report.html

```
from ydata_profiling import ProfileReport
```

```
# Generate the EDA report
eda_report = ProfileReport(df_filtered, title="Grain Export EDA Report", explorative=True)

# Save the report as an HTML file
# Use a valid path where you want to save the report
# For example, to save it in the current directory, use:
eda_report_path = "EDA_Report.html"
eda_report.to_file(eda_report_path)

# Return the file path for download (if needed)
eda_report_path
```

```

/usr/local/lib/python3.11/dist-packages/ydata_profiling/utils/dataframe.py:137: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view
df.rename(columns={"index": "df_index"}, inplace=True)

Summarize dataset: 100%
24/04 10:02:00.00 5.6414s Completed

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load the merged and cleaned dataset assumed to be the same used in EDA
# For visualization, let's simulate a sample version of it since actual data was not provided
# Example structure: crop_year, grain_type, region, grain_week, kt

# Simulated example data for visualization
data = {
    "grain_week": list(range(1, 53)) * 2,
    "kt": [abs(1000 + 200 * (i % 10) + (i % 5) * 100) for i in range(52)] +
          [abs(1200 + 150 * (i % 8) + (i % 3) * 200) for i in range(52)],
    "grain_type": ["Corn"] * 52 + ["Soybeans"] * 52,
    "region": ["Bay & Lakes"] * 26 + ["St. Lawrence"] * 26 + ["Bay & Lakes"] * 26 + ["St. Lawrence"] * 26
}

df = pd.DataFrame(data)

# Plot 1: Seasonality across grain weeks
plt.figure(figsize=(12, 6))
sns.lineplot(data=df, x="grain_week", y="kt", hue="grain_type", style="region", markers=True)
plt.title("Weekly Grain Export Trends by Type and Region")
plt.xlabel("Grain Week")
plt.ylabel("Export Volume (kt)")
plt.grid(True)
plt.tight_layout()
plt.show()

# Plot 2: Export Volume Distribution
plt.figure(figsize=(10, 6))
sns.boxplot(data=df, x="grain_type", y="kt", hue="region")
plt.title("Export Volume Distribution by Grain Type and Region")
plt.ylabel("Export Volume (kt)")
plt.grid(True)
plt.tight_layout()
plt.show()

# Plot 3: Heatmap of average exports by week and region
pivot = df.pivot_table(values='kt', index='grain_week', columns='region', aggfunc='mean')
plt.figure(figsize=(12, 6))
sns.heatmap(pivot, cmap="YlGnBu", annot=False, linewidths=0.5)

```